

Lecture 15

Parallel I/O

Overview

- What is parallel I/O
- MPI-I/O
- Parallel File Systems
- Application examples

Parallel I/O

- I/O for scientific computing
 - Not for databases, enterprise applications, the Internet
- Large data items
 - Arrays
 - Simulation results
- Shared across many processors
 - Hundreds or thousands

An Example Application

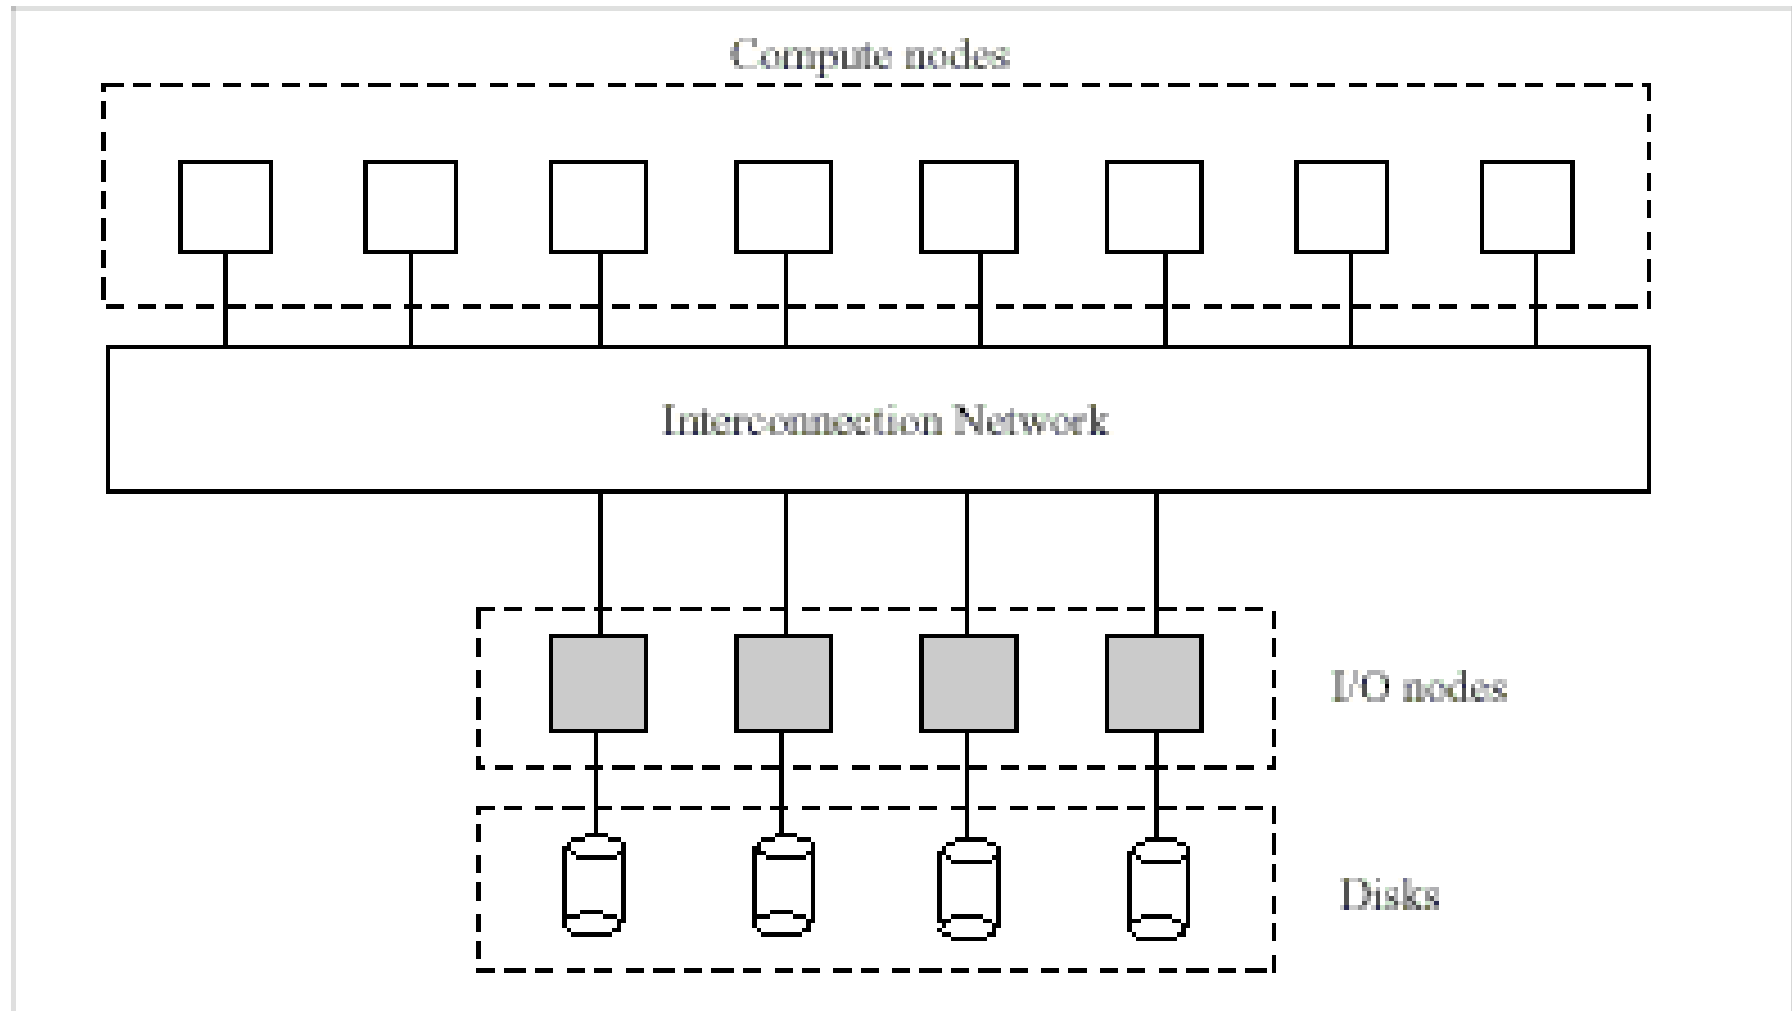
The I/O Gap in Parallel Applications

- The I/O gap between memory speed and average disk access stands at 10^5 (40 nanosecs versus 4 ms)
- Is the gap larger or smaller for parallel applications and why?

The I/O Gap in Parallel Applications

- The I/O gap between memory speed and average disk access stands at 10^5 (40 nanosecs versus 4 ms)
- Is the gap larger or smaller for parallel applications and why?
 - Better: caches are larger and shared through distributed memory systems?
 - Worse: many access patterns are sequential and therefore unshareable?
 - Worse: scientific applications tend to be write-oriented

Parallel I/O Architecture




Parallel I/O Architecture

- I/O nodes
 - Block servers
 - Implement availability and caching (RAID)
 - Usually dedicated (no computation)
- Processors nodes
 - Sometimes have their own disk (not part of the compute system)
- High speed interconnect
 - Often switches
- Example systems
 - IBM SP/2 (ASCI White, BME has one)
 - Intel Tflops (ASCI Red)
 - Cray T3E
 - IBM BlueGene/L

The API Problem


- UNIX has a “standard” API
 - `open()`, `close()`, `read()`, `write()`
- Parallel file systems often want to access:
 - Small data items from within a file
 - Non-contiguous in the file space
- Example
 - Two-dimensional array (spatial data)
 - Stored in row-major order
 - Partition (to processors) based on rectangles

Typical Array Workload



Typical Array Workload

- UNIX interface requires seeks between each I/O
 - Interactive/blocking interface
- No way to encode, overall view of access pattern
 - Application has a priori knowledge of data usage
- Workload is not aligned to file layout (sequential)
 - Generally difficult to lay out sequentially, because used in different ways at different times
- Not interactive
 - Latency insensitive



The API Problem (again)

- UNIX has a “standard” API
 - `open()`, `close()`, `read()`, `write()`
- Parallel file systems often want to access:
 - Small data items from within a file
 - Non-contiguous in the file space
- Example
 - Two-dimensional array (spatial data)
 - Stored in row-major order
 - Partition (to processors) based on rectangles
- How do applications perform parallel I/O in a standard, portable way?
 - Well they don't really.

MPI-I/O

- An API standard for parallel I/O
 - Much different than POSIX/UNIX API
- But, as it gains traction, becomes easier to use
- Replaces many custom libraries and approaches
 - PANDA
 - Solar
 - ChemI/O
 - Others ...

MPI-I/O History


- Which of the following are true?
 - MPI-I/O grew out of a study that looked at how message passing architectures and standards affected I/O
 - MPI-I/O can be used for both good and evil
 - MPI-I/O was started at IBM Watson
 - MPI-I/O is easy to use
 - MPI-I/O maps I/O reads and writes to message passing sends and receives

MPI-I/O History

- Which of the following are true?
 - MPI-I/O grew out of a study that looked at how message passing architectures and standards affected I/O
 - MPI-I/O can be used for both good and evil
 - MPI-I/O was started at IBM Watson
 - MPI-I/O is easy to use
 - MPI-I/O maps I/O reads and writes to message passing sends and receives

MPI-I/O for Dum-Dummies

- If, n processors wish to read a file in n chunks
- `MPI_File_Open` – does an open for all n processors and indicates a “communicator” that defines the group
- Each processor does `MPI_File_seek`, `MPI_File_Read`, and `MPI_File_Write`
 - *i.e.*, each processor has its own seek pointer
- `MPI_File_Read/Write` operate on typed arrays of data
 - Meaningless semantic sugar



MPI-I/O for Jane and John Smith


- Multiple I/O interfaces
 - Implicit offset (seek pointer)
 - Explicit offset
 - Shared file pointer (among all processes)
 - *E.g.*, a bunch of processes that read the next data item when completing work against a previous data item
- Multiple storage representations (recall I/Os are typed)
 - Native (no memory interpretation)
 - Internal (app provides interface for how to convert memory to persistent data)
 - External32 (network representation of data, similar to TCP)

MPI-I/O for Savants

- Many optimizations for groups of processes to do I/O
 - Collective I/O
 - Data sieving
- Hints
 - Information to help applications choose striping, prefetching, read-ahead and caching policies

Data Sieving

- Package a series of non-contiguous requests into a single large read
 - Use sequential capabilities of the underlying devices/systems
- Make as few requests as possible and make them big

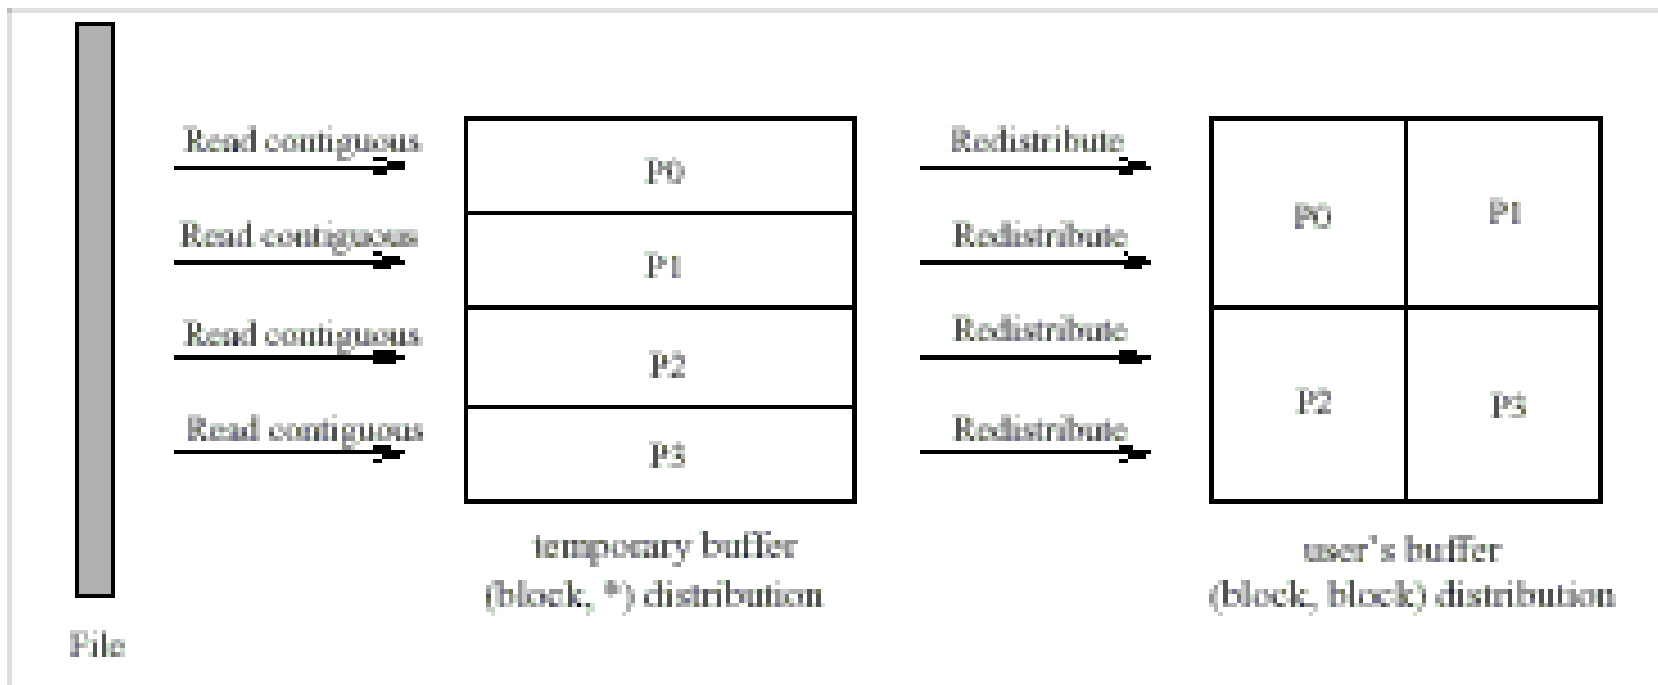


Collective I/O

- Create an I/O plan across many processors
 - Generally some communication to set-up the task
- Main technique – two-phase I/O
- Other techniques
 - Disk and server directed I/O
 - Similar, but limit the buffering requirements at I/O nodes

Two-Phase I/O

- I/Os are done in chunks consistent with the disk layout
- I/Os are redistributed into users buffers according to the application (in-memory) view of data
 - Redistributed among processors
- Same array example



Collective I/O




Figure 2: A file access by four processors.




Figure 3: File access using the collective I/O.

How good is MPI I/O?

- Depends on
 - How it is used?
 - What the implementation does with the information it has?
- MPI I/O is an interface
 - Many implementations are possible
 - Popular implementation is ROMIO, from the guys who make the graphs/figures

Level #0 – Ogres Have Layers

- Many independent, contiguous requests
 - No access information available to MPI system at runtime

```
MPI_File_Open ( ..., "filename", ..., &fh );  
for ( i=0; i<n_rows; i++ )  
{  
    MPI_File_Seek ( fh, ... );  
    MPI_File_read ( fh, row[i], ... );  
}
```


Level #1 – Ogres Have Layers

- Many collective, contiguous requests
 - MPI implementation expects to see same access pattern at multiple sites
 - Can lead to good read-ahead, prefetching decisions when implementation sees patterns repeat at different processors

```
MPI_File_Open ( MPI_COMM_WORLD, "filename", ..., &fh );  
for ( i=0; i<n_rows; i++ )  
{  
    MPI_File_Seek ( fh, ... );  
    MPI_File_read ( fh, row[i], ... );  
}
```

Level #2 – Ogres Have Layers

- Single independent, non-contiguous requests
 - Data sieving can be used
 - Based on an application defined data type

`MPI_Type_create_subarray (... , &subarray,);`

`MPI_Type_commit (&subarray);`

`MPI_File_open (..., "filename", ..., &fh);`

`MPI_File_set_view (fh, ..., &subarray,);`

`MPI_File_read (fh, local_array, ...);`

Level #3 – Ogres Have Layers

- Multiple collective, non-contiguous requests
 - Data sieving can be used
 - Collective I/O can be used

```
MPI_Type_create_subarray (... , &subarray, .... );
```

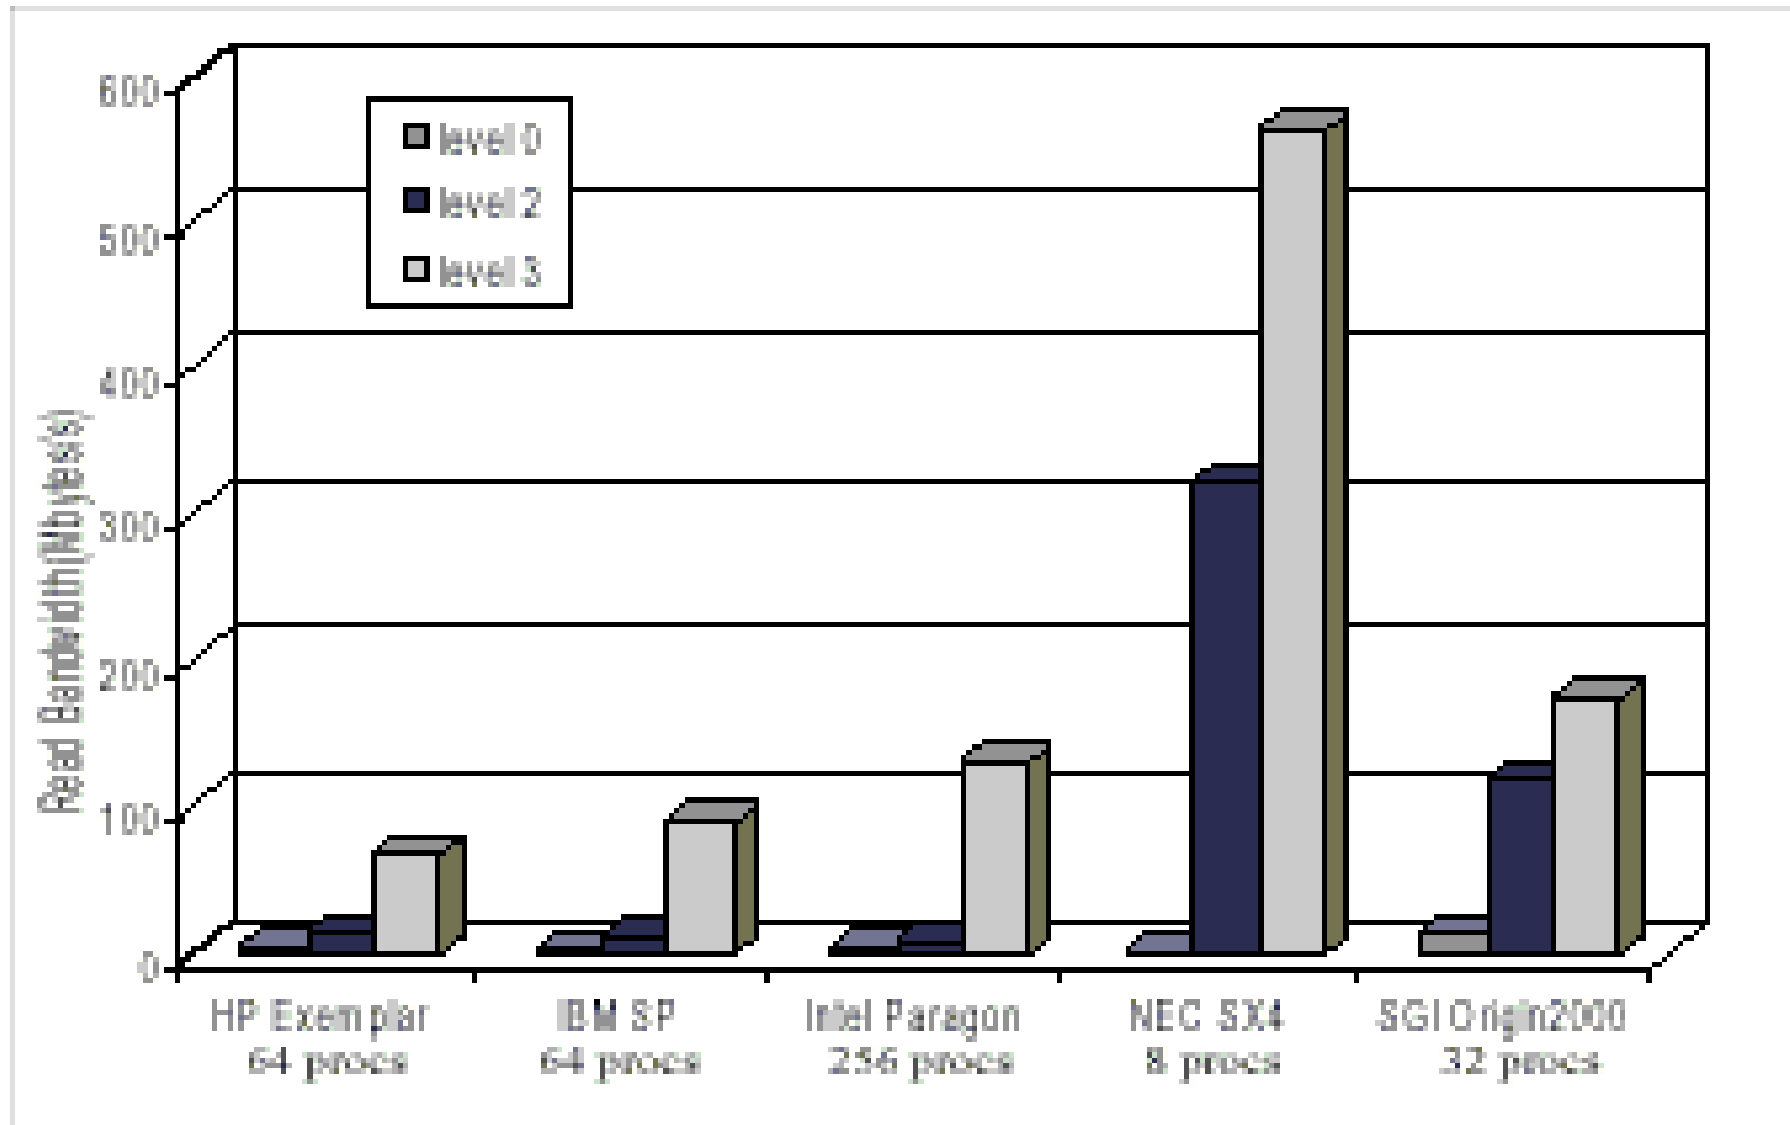
```
MPI_Type_commit ( &subarray );
```

```
MPI_File_open ( MPI_COMM_WORLD, "filename", ..., &fh );
```


```
MPI_File_set_view ( fh, ..., &subarray, .... );
```

```
MPI_File_read ( fh, local_array, ... );
```

MPI I/O Performance (Read)



MPI I/O Performance (Write)



Performance Observations

- Sieving (level #2) performs well on read
- Write requires collective I/O to get high bandwidth
- Scaling is highly dependent on system architecture


Is MPI-I/O Important

- The developers of GPFS tell me that their users almost always use the POSIX interface, not MPI-I/O
- Can storage system technologies mimic the function of MPI-I/O implicitly?
 - Data sieving?
 - Collective I/O?


SANs and High-Perf. Computing

- FC/iSCSI SANs create block storage endpoints
 - Similar to the storage servers in the HP model
- Can SAN storage replace storage servers
 - In GPFS for the POSIX interface?
 - When using MPI-I/O?


Multi-Collective I/O



Multi-Collective I/O



Multi-Collective I/O



Multi-Collective I/O

