

MPI-2 Features

MPI-2 Features Overview

- **Parallel I/O** specified
- Dynamic **process creation**
- Extended collective operations
- One-sided communications
- Bindings for C++ (big disappointment)
- Bindings for Fortran 2000+ (smaller disappointment)
- Thread information access

MPI-2 and Threads

- Threads can be used for parallelism and concurrency (typically, waiting for events)
- MPI **allows** thread safety, **does not require** it
 - MPI-2 supplies inquiry functions for level of thread safety in the MPI implementation
 - Only MPI_Probe() involves testing for "current state" information in MPI-1 (*Probe waits until a message matching source, tag, and comm arrives*)

MPI-2 and Threads

- `MPI_Init_thread(int *argc, char **argv, int *provided)`
- Additional argument provided requests, returns level
 - `MP_THREAD_SINGLE`: only one user thread
 - `MP_THREAD_FUNNELED`: multiple threads allowed, but only one user thread makes MPI calls
 - `MP_THREAD_SERIALIZED`: mutiple threads allowed, but only one at a time makes MPI calls
 - `MP_THREAD_MULTIPLE`: all allowed

MPI-2 and Threads

- Thread that calls `MPI_Init_thread` is main thread
 - Only one allowed to call `MPI_Finalize`
- Value of requested thread level need not be same on all processes calling `MPI_Init_thread`
- Thread creation, management is not part of MPI
 - left to existing thread packages
 - using `pthread`s, `OpenMP` as shown

MPI-2 Externals

- External interface specs for debuggers, other tools to attach to MPI (or to layer MPI-2 on top of MPI-1)
 - `MPI_Type_get_envelope/contents` for decoding an MPI derived datatype
 - `Generalized requests`: user can define new nonblocking messaging functions, then use `MPI_Test`, `MPI_Wait`, etc.
- Explicit handling of C/Fortran interoperability
 - E.g., type `MPI_Fint` in C

MPI-2 Odds and Ends

- Can use collective communications with *intercommunicators* as well as *intracommunicators*
 - MPI_AllReduce, MPI_Bcast most common
- Intracommunicator collective calls can have an **in-place** option, allowing send and receive buffers to be same.
 - Recall, aliasing explicitly not allowed in MPI !
- **Attribute caching** for more than just communicators
 - datatypes, windows (used for RMA)

MPI-2 Odds and Ends

- One-sided communications (remote memory ops)
 - Looks like *shmem* programs; one process can essentially send w/o other explicitly receiving
 - Based on idea of **window**: portion of a process's address space it offers for remote memory ops by other MPI processes
 - Use **put**, **get**, **accumulate** to store to, load from, add into another process's window
 - Not implemented by some MPI implementations; buggy, unoptimized

MPI2: Process Creation

- MPI-1 is static: no processes after startup
- **PVM** was built around same time as MPI
 - parallel virtual machine
 - provides some distributed OS facilities
 - processes can be created, destroyed dynamically
 - can access information via *mailboxes*
 - based at Oak Ridge National Lab, still actively used
 - part of most Linux distributions (like OpenMPI)

MPI2: Process Management

- MPI-1 is strictly an *interface*, avoids OS jobs like resource management
- MPI-2 process mgmt operations are collective among all processes (ones creating, and ones being created)
- Resulting sets are represented by an intercomm
- Two fundamental operations
 - spawning new tasks
 - connecting to them

MPI2: Process Creation

- `MPI_COMM Spawn` starts MPI processes and establishes communication with them, returning an intercommunicator.
- `MPI_COMM Spawn Multiple` starts several different binaries (or same binary with different arguments).
- `MPI_Info()` is a hook to allow communication with underlying mechanism for distributed OS.

MPI2: Process Creation

- Attribute `MPI_UNIVERSE_SIZE` of `MPI_COMM_WORLD` gives total number of processes can usefully be started in all.
- `UNIVERSE - WORLD` is how many processes might usefully be started in addition to those already running.
- Note: implementation can return a value that is not anywhere near useful.

New Process Communication

- MPI-2 can establish communication between two sets of MPI processes that *do not* share a communicator
- Collective but asymmetric process: one group is server, another is client.
- How does client group discover server group?
 - MPI_Open_port + MPI_Comm_accept, or
 - MPI_Publish_name + MPI_Lookup_name

MPI C++ bindings

- MPI was intended partly be multilanguage
 - Call a C library from Fortran, vice-versa
 - C++ did not exist at time (but Fortran 90 did)
 - Lead to "lowest common denominator"
- What object should a base C++ class for MPI represent?
- No direct way to pass objects, just pack/unpack their data members

MPI C++ bindings

- Thin layer on top of MPI-C, wrapping mpi.h in an *extern "C" {...}*
- Names of member functions are "consistent with" the underlying MPI functions.
- Uses namespaces

C++ Bindings

```
namespace MPI {  
    class Comm                {...};  
    class Intracomm : public Comm    {...};  
    class Graphcomm : public Intracomm    {...};  
    class Cartcomm : public Intracomm    {...};  
    class Intercomm : public Comm    {...};  
    class Datatype            {...};  
    class Errhandler          {...};  
    class Exception           {...};  
    class Group               {...};  
    class Op                  {...};  
    class Request              {...};  
    class Prerequest : public Request    {...};  
    class Status               {...};  
    class File                 {...};  
    class Grequest : public Request    {...};  
    class Info                 {...};  
    class Win                  {...}; }  
}
```


roundrobin with C++-bindings

```
#include <mpi++.h> ... // Note I use comm_world here but should use different name
MPI::Init(argc, argv);
int rank = MPI::COMM_WORLD.Get_rank();
int size = MPI::COMM_WORLD.Get_size();
int to    = (rank + 1) % p;
int from = (p + rank - 1) % p;
if (rank == p - 1)
    MPI::COMM_WORLD.Send(&msg, 1, MPI::INT, to, 4);
MPI::COMM_WORLD.Recv(&msg, 1, MPI::INT, from, MPI::ANY_TAG);
MPI::COMM_WORLD.Send(&msg, 1, MPI::INT, to, 4);
if (rank == 0) {
    MPI::COMM_WORLD.Recv(&msg, 1, MPI::INT, from, MPI::ANY_TAG);
    cout << "Node " << rank << " received " << msg << endl; }
MPI::Finalize(); ...}
```

OOMPI

- Library specification providing higher level OO
- A *port* is an object encapsulating *src* and *dest* of a message (a genuine connection model)
- Has info about underlying MPI communicator and the rank of destination
- i^{th} port of a communicator is array-indexed

```
OOPMPI_Intra_comm rowcomm;  
rowcomm[m].Send(i);    // send i from port m  
rowcomm[n].Receive(j); // recv j from port n
```

OOMPI Messages

- OOMPI_Message object supports basic datatypes
- OOMPI_Array_message for arrays
- User-defined types, packing similar to C versions
- Avoid explicit message construction when possible

```
OOPMPI_Port port;  
double v[100]; int tag;  
...  
port.Send(v, 100, tag)
```

OOMPI Messages

- OOMPI_Port::Send(OOMPI_Message buffer, int tag = OOMPI_NO_TAG);
- OOMPI_Port::Send(OOMPI_Array_message buffer, int count int tag = OOMPI_NO_TAG);
- Allows users to avoid explicit tagging if desired
- Stream interface also allowed (but inefficient)

```
OOMPI_Port port;  
...  
port << i << j;  
port >> i >> j;
```

roundrobin in OOPMPI

```
OOMPI_COMM_WORLD.Init(argc, argv);
int rank = OOMPI_COMM_WORLD.Rank( );
int p = OOMPI_COMM_WORLD.Size( );
OOMPI_Port to = OOMPI_COMM_WORLD[(rank+1)%p];
OOMPI_Port from = OOMPI_COMM_WORLD[(p+rank-1)%p];

if (rank == p-1) to << msg;
from >> msg;
to << msg;
if (rank == 0) from >> msg;

OOMPI_COMM_WORLD.Finalize( );
```